

# FLEX SCHEMAS WITH POSTGRESQL AND DJANGO

Advanced PostgreSQL: Models halfway between Relational and NoSQL

---

José Antonio Perdiguero López

 <https://github.com/PeRDy>

 <https://www.linkedin.com/in/josé-antonio-perdiguero-lópez-732a9768/>

@ perdy.hh@gmail.com

May 5, 2017

Málaga Python MeetUp

OpenSouthCode 2017

Flex schemas with PostgreSQL and Django

2017-05-05

FLEX SCHEMAS WITH POSTGRESQL AND DJANGO

Advanced PostgreSQL: Models halfway between Relational and NoSQL

---

José Antonio Perdiguero López

 <https://github.com/PeRDy>

 <https://www.linkedin.com/in/josé-antonio-perdiguero-lópez-732a9768/>

 perdyhh@gmail.com

May 5, 2017

Málaga Python MeetUp

OpenSouthCode 2017

1	Introduction	3	Indexes
	Relational Databases		Hash
	NoSQL Databases		B-Tree
	Relational vs NoSQL		GiST
2	Fields		SP-GiST
	ArrayField		GIN
	HStoreField		BRIN
	JSONField	4	Conclusion
			PostgreSQL as Relational and
			NoSQL
			Use cases

2017-05-05



Index

1	Introduction	3	Indexes
	Relational Databases		Hash
	NoSQL Databases		B-Tree
	Relational vs NoSQL		GiST
2	Fields		SP-GiST
	ArrayField		GIN
	HStoreField		BRIN
	JSONField	4	Conclusion
			PostgreSQL as Relational and
			NoSQL
			Use cases

## Introduction

---

2017-05-05

Flex schemas with PostgreSQL and Django

└ Introduction

Introduction

---

Relational databases are those based on *Relational Model*, invented by *E.F. Codd*.

The relational model define that all the data is represented in terms of **tuples**, grouped in **relations**. These relations consists of a set of attributes and a set of  $n$  tuples, and are called **tables**.

2017-05-05

Flex schemas with PostgreSQL and Django  
└─ Introduction  
    └─ Relational Databases  
        └─ Relational Databases

Relational databases are those based on *Relational Model*, invented by *E.F. Codd*.

The relational model define that all the data is represented in terms of **tuples**, grouped in **relations**. These relations consists of a set of attributes and a set of  $n$  tuples, and are called **tables**.



- Structured data
- Strong typing
- Integrity and constraints
- Indexes
- Support
- Fast querying
- Maturity



- Admin
- Poor horizontal scaling
- Schema and data migrations

2017-05-05

Flex schemas with PostgreSQL and Django  
└─ Introduction  
    └─ Relational Databases  
        └─ Relational Databases: Pros vs Cons



- Structured data
- Strong typing
- Integrity and constraints
- Indexes
- Support
- Fast querying
- Maturity



- Admin
- Poor horizontal scaling
- Schema and data migrations

## Key-Value

Key-Value stores are the simplest NoSQL data stores. The value is a blob, and due to primary-key access they have a great performance and scalability.

## Document

Documents are the main concept in document databases. The database stores and retrieves documents, which can be XML, JSON, BSON, and so on.

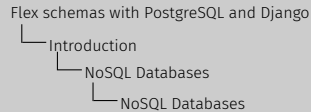
## Column family

Column-family databases store data in column families as rows that have many columns associated with a row key.

## Graph

Graph databases allow you to store entities and relationships between these entities. Entities are also known as nodes, which have properties. Relations are known as edges that can have properties.

2017-05-05



### Key-Value

Key-Value stores are the simplest NoSQL data stores. The value is a blob, and due to primary-key access they have a great performance and scalability.

### Document

Documents are the main concept in document databases. The database stores and retrieves documents, which can be XML, JSON, BSON, and so on.

### Column family

Column-family databases store data in column families as rows that have many columns associated with a row key.

### Graph

Graph databases allow you to store entities and relationships between these entities. Entities are also known as nodes, which have properties. Relations are known as edges that can have properties.

Memcached and Redis as **Key-Value** databases.

MongoDB as **document** database.

Cassandra is a **column family** database. These databases are like relational tables but each row can have a different subset of columns.

In **graph** databases think of a node as an instance of an object in the application. OrientDB and Neo4j.



- Fast writing
- Horizontal scaling
- Reduced admin tasks
- Flexible
- Mostly open-source



- Slow querying
- Support
- Maturity

2017-05-05

Flex schemas with PostgreSQL and Django  
└─ Introduction  
    └─ NoSQL Databases  
        └─ NoSQL Databases: Pros vs Cons



- Fast writing
- Horizontal scaling
- Reduced admin tasks
- Flexible
- Mostly open-source



- Slow querying
- Support
- Maturity

# Should I use a Relational or NoSQL database?



Flex schemas with PostgreSQL and Django

└─ Introduction

└─ Relational vs NoSQL

2017-05-05

Should I use a Relational or NoSQL database?





## Fields

---

2017-05-05

Flex schemas with PostgreSQL and Django

└ Fields

Fields

---

A field for storing lists of data that can be nested to store multi-dimensional arrays.

## Example

```
from django.db import models
from django.contrib.postgres.fields import ArrayField

class Post(models.Model):
    name = models.CharField(max_length=200)
    bar = ArrayField(models.CharField(max_length=200), blank=True)

    def __str__(self):
        return self.name

>>> Post.objects.create(name='First post', tags=['thoughts', 'django'])
>>> Post.objects.create(name='Second post', tags=['thoughts'])
>>> Post.objects.create(name='Third post', tags=['tutorial', 'django'])
```

2017-05-05

Flex schemas with PostgreSQL and Django



A field for storing lists of data that can be nested to store multi-dimensional arrays.

### Example

```
from django.db import models
from django.contrib.postgres.fields import ArrayField

class Post(models.Model):
    name = models.CharField(max_length=200)
    bar = ArrayField(models.CharField(max_length=200), blank=True)

    def __str__(self):
        return self.name

>>> Post.objects.create(name='First post', tags=['thoughts', 'django'])
>>> Post.objects.create(name='Second post', tags=['thoughts'])
>>> Post.objects.create(name='Third post', tags=['tutorial', 'django'])
```

## Contains (@>)

```
>>> Post.objects.filter(tags__contains=['thoughts'])
<QuerySet [<Post: First post>, <Post: Second post>]>
```

```
>>> Post.objects.filter(tags__contains=['django'])
<QuerySet [<Post: First post>, <Post: Third post>]>
```

```
>>> Post.objects.filter(tags__contains=['django', 'thoughts'])
<QuerySet [<Post: First post>]>
```

## Contained By (<@)

```
>>> Post.objects.filter(tags__contained_by=['thoughts', 'django'])
<QuerySet [<Post: First post>, <Post: Second post>]>
```

```
>>> Post.objects.filter(tags__contained_by=['thoughts', 'django', 'tutorial'])
<QuerySet [<Post: First post>, <Post: Second post>, <Post: Third post>]>
```

2017-05-05

Flex schemas with PostgreSQL and Django

Fields

ArrayField

ArrayField: Queries

### Contains (@>)

```
>>> Post.objects.filter(tags__contains=['thoughts'])
<QuerySet [<Post: First post>, <Post: Second post>]>
```

```
>>> Post.objects.filter(tags__contains=['django'])
<QuerySet [<Post: First post>, <Post: Third post>]>
```

```
>>> Post.objects.filter(tags__contains=['django', 'thoughts'])
<QuerySet [<Post: First post>]>
```

### Contained By (<@)

```
>>> Post.objects.filter(tags__contained_by=['thoughts', 'django'])
<QuerySet [<Post: First post>, <Post: Second post>]>
```

```
>>> Post.objects.filter(tags__contained_by=['thoughts', 'django', 'tutorial'])
<QuerySet [<Post: First post>, <Post: Second post>, <Post: Third post>]>
```

**First post:** thoughts, django.

**Second post:** thoughts.

**Third post:** tutorial, django.

Remember here the relevance of the lookups and transform defined by Django, because these will be the most widely used, so it's a good idea to **index** them.

## Overlap (&&)

```
>>> Post.objects.filter(tags__overlap=['thoughts'])
<QuerySet [<Post: First post>, <Post: Second post>]>
```

```
>>> Post.objects.filter(tags__overlap=['thoughts', 'tutorial'])
<QuerySet [<Post: First post>, <Post: Second post>, <Post: Third post>]>
```

## Len

```
>>> Post.objects.filter(tags__len=1)
<QuerySet [<Post: Second post>]>
```

2017-05-05

Flex schemas with PostgreSQL and Django

Fields

ArrayField

ArrayField: Queries

### Overlap (&&)

```
>>> Post.objects.filter(tags__overlap=['thoughts'])
<QuerySet [<Post: First post>, <Post: Second post>]>

>>> Post.objects.filter(tags__overlap=['thoughts', 'tutorial'])
<QuerySet [<Post: First post>, <Post: Second post>, <Post: Third post>]>
```

### Len

```
>>> Post.objects.filter(tags__len=1)
<QuerySet [<Post: Second post>]>
```

**First post:** thoughts, django.

**Second post:** thoughts.

**Third post:** tutorial, django.

Remember here the relevance of the lookups and transform defined by Django, because these will be the most widely used, so it's a good idea to **index** them.

## Index Transforms

```
>>> Post.objects.filter(tags__0='thoughts')
<QuerySet [<Post: First post>, <Post: Second post>]>
```

```
>>> Post.objects.filter(tags__1__iexact='Django')
<QuerySet [<Post: First post>]>
```

```
>>> Post.objects.filter(tags__276='javascript')
<QuerySet []>
```

## Slice Transforms

```
>>> Post.objects.filter(tags__0_1=['thoughts'])
<QuerySet [<Post: First post>, <Post: Second post>]>
```

```
>>> Post.objects.filter(tags__0_2__contains=['thoughts'])
<QuerySet [<Post: First post>, <Post: Second post>]>
```

2017-05-05

Flex schemas with PostgreSQL and Django

Fields

ArrayField

ArrayField: Queries

### Index Transforms

```
>>> Post.objects.filter(tags__0='thoughts')
<QuerySet [<Post: First post>, <Post: Second post>]>
```

```
>>> Post.objects.filter(tags__1__iexact='Django')
<QuerySet [<Post: First post>]>
```

```
>>> Post.objects.filter(tags__276='javascript')
<QuerySet []>
```

### Slice Transforms

```
>>> Post.objects.filter(tags__0_1=['thoughts'])
<QuerySet [<Post: First post>, <Post: Second post>]>
```

```
>>> Post.objects.filter(tags__0_2__contains=['thoughts'])
<QuerySet [<Post: First post>, <Post: Second post>]>
```

**First post:** thoughts, django.

**Second post:** thoughts.

**Third post:** tutorial, django.

Remember here the relevance of the lookups and transform defined by Django, because these will be the most widely used, so it's a good idea to **index** them.

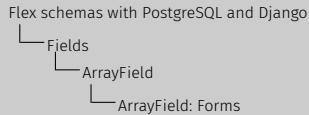
## SimpleArrayField

A simple field which maps to an array. It is represented by an HTML `<input>`.

## SplitArrayField

This field handles arrays by reproducing the underlying field a fixed number of times.

2017-05-05



### SimpleArrayField

A simple field which maps to an array. It is represented by an HTML `<input>`.

### SplitArrayField

This field handles arrays by reproducing the underlying field a fixed number of times.

A field for storing key-value pairs. The Python data type used is a dict that maps *strings* into *nullable strings*.

## Example

```
from django.contrib.postgres.fields import HStoreField
from django.db import models

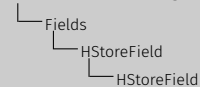
class Dog(models.Model):
    name = models.CharField(max_length=200)
    data = HStoreField()

    def __str__(self):
        return self.name

>>> Dog.objects.create(name='Rufus',
                        data={'breed': 'labrador', 'owner': 'Bob', 'friend': 'Bobby'})
>>> Dog.objects.create(name='Meg',
                        data={'breed': 'collie', 'owner': 'Bob', 'toy': 'yellow ball'})
>>> Dog.objects.create(name='Fred', data={})
```

Flex schemas with PostgreSQL and Django

2017-05-05



A field for storing key-value pairs. The Python data type used is a dict that maps *strings* into *nullable strings*.

### Example

```
from django.contrib.postgres.fields import HStoreField
from django.db import models

class Dog(models.Model):
    name = models.CharField(max_length=200)
    data = HStoreField()

    def __str__(self):
        return self.name

>>> Dog.objects.create(name='Rufus',
                        data={'breed': 'labrador', 'owner': 'Bob', 'friend': 'Bobby'})
>>> Dog.objects.create(name='Meg',
                        data={'breed': 'collie', 'owner': 'Bob', 'toy': 'yellow ball'})
>>> Dog.objects.create(name='Fred', data={})
```

Remember that for using HStoreField is required to install **hstore extension**.

A field for storing key-value pairs. The Python data type used is a dict that maps *strings* into *nullable strings*.

## PostgreSQL extension

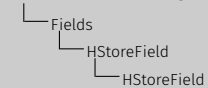
```
from django.contrib.postgres.operations import HStoreExtension

class Migration(migrations.Migration):
    ...

    operations = [
        HStoreExtension(),
        ...
    ]
```

Flex schemas with PostgreSQL and Django

2017-05-05



A field for storing key-value pairs. The Python data type used is a dict that maps *strings* into *nullable strings*.

### PostgreSQL extension

```
from django.contrib.postgres.operations import HStoreExtension

class Migration(migrations.Migration):
    ...

    operations = [
        HStoreExtension(),
        ...
    ]
```

Remember that for using HStoreField is required to install **hstore extension**.



Key Lookups

```
>>> Dog.objects.filter(data__breed='collie')
<QuerySet [<Dog: Meg>]>

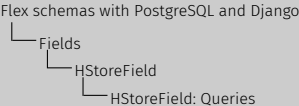
>>> Dog.objects.filter(data__breed__contains='l')
<QuerySet [<Dog: Rufus>, <Dog: Meg>]>
```

Contains (@>)

```
>>> Dog.objects.filter(data__contains={'owner': 'Bob'})
<QuerySet [<Dog: Rufus>, <Dog: Meg>]>

>>> Dog.objects.filter(data__contains={'breed': 'collie'})
<QuerySet [<Dog: Meg>]>
```

2017-05-05



Key Lookups

```
>>> Dog.objects.filter(data__breed='collie')
<QuerySet [<Dog: Meg>]>

>>> Dog.objects.filter(data__breed__contains='l')
<QuerySet [<Dog: Rufus>, <Dog: Meg>]>
```

Contains (@>)

```
>>> Dog.objects.filter(data__contains={'owner': 'Bob'})
<QuerySet [<Dog: Rufus>, <Dog: Meg>]>

>>> Dog.objects.filter(data__contains={'breed': 'collie'})
<QuerySet [<Dog: Meg>]>
```

Rufus:	breed: labrador. owner: Bob. friend: Bobby.
Meg:	breed: collie. owner: Bob. toy: yellow ball.
Fred:	

Keys example uses **overlap** because order is not guaranteed.

## Contained By (<@)

```
>>> Dog.objects.filter(data__contained_by={'breed': 'collie', 'owner': 'Bob'})
<QuerySet [<Dog: Meg>, <Dog: Fred>]>
```

```
>>> Dog.objects.filter(data__contained_by={'breed': 'collie'})
<QuerySet [<Dog: Fred>]>
```

## Has Key (?)

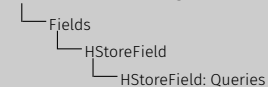
```
>>> Dog.objects.filter(data__has_key='toy')
<QuerySet [<Dog: Meg>]>
```

## Has Any Keys (?|)

```
>>> Dog.objects.filter(data__has_any_keys=['toy', 'friend'])
<QuerySet [<Dog: Rufus>, <Dog: Meg>]>
```

Flex schemas with PostgreSQL and Django

2017-05-05



### Contained By (<@)

```
>>> Dog.objects.filter(data__contained_by={'breed': 'collie', 'owner': 'Bob'})
<QuerySet [<Dog: Meg>, <Dog: Fred>]>
```

```
>>> Dog.objects.filter(data__contained_by={'breed': 'collie'})
<QuerySet [<Dog: Fred>]>
```

### Has Key (?)

```
>>> Dog.objects.filter(data__has_key='toy')
<QuerySet [<Dog: Meg>]>
```

### Has Any Keys (?|)

```
>>> Dog.objects.filter(data__has_any_keys=['toy', 'friend'])
<QuerySet [<Dog: Rufus>, <Dog: Meg>]>
```

Rufus:	breed: labrador. owner: Bob. friend: Bobby.
Meg:	breed: collie. owner: Bob. toy: yellow ball.
Fred:	

Keys example uses **overlap** because order is not guaranteed.

Has Keys (?&)

```
>>> Dog.objects.filter(data__has_keys=['breed', 'owner'])
<QuerySet [<Dog: Rufus>, <Dog: Meg>]>
```

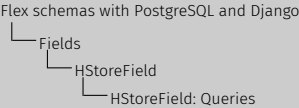
Keys

```
>>> Dog.objects.filter(data__keys__overlap=['breed', 'toy'])
<QuerySet [<Dog: Rufus>, <Dog: Meg>]>
```

Values

```
>>> Dog.objects.filter(data__values__contains=['collie'])
<QuerySet [<Dog: Meg>]>
```

2017-05-05



Has Keys (?&)

```
>>> Dog.objects.filter(data__has_keys=['breed', 'owner'])
<QuerySet [<Dog: Rufus>, <Dog: Meg>]>
```

Keys

```
>>> Dog.objects.filter(data__keys__overlap=['breed', 'toy'])
<QuerySet [<Dog: Rufus>, <Dog: Meg>]>
```

Values

```
>>> Dog.objects.filter(data__values__contains=['collie'])
<QuerySet [<Dog: Meg>]>
```

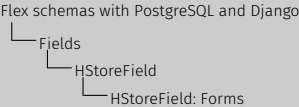
Rufus:	breed: labrador. owner: Bob. friend: Bobby.
Meg:	breed: collie. owner: Bob. toy: yellow ball.
Fred:	

Keys example uses **overlap** because order is not guaranteed.

HStoreField

A field which accepts JSON encoded data for an HStoreField, casting all values (except nulls) to strings. It is represented by an HTML `<textarea>`.

2017-05-05



**HStoreField**  
A field which accepts JSON encoded data for an HStoreField, casting all values (except nulls) to strings. It is represented by an HTML `<textarea>`.

# JSONField

A field for storing JSON encoded data. In Python the data is represented in its Python native format: dictionaries, lists, strings, numbers, booleans and None.

## Example

```
from django.contrib.postgres.fields import JSONField
from django.db import models
```

```
class Dog(models.Model):
    name = models.CharField(max_length=200)
    data = JSONField()
```

```
    def __str__(self):
        return self.name
```

```
>>> Dog.objects.create(name='Rufus', data={
...     'breed': 'labrador',
...     'owner': {
...         'name': 'Bob',
...         'other_pets': [{'name': 'Fishy'}],
...     },
... })
>>> Dog.objects.create(name='Meg', data={'breed': 'collie'})
```

Flex schemas with PostgreSQL and Django

2017-05-05



## JSONField

A field for storing JSON encoded data. In Python the data is represented in its Python native format: dictionaries, lists, strings, numbers, booleans and None.

### Example

```
from django.contrib.postgres.fields import JSONField
from django.db import models

class Dog(models.Model):
    name = models.CharField(max_length=200)
    data = JSONField()

    def __str__(self):
        return self.name

>>> Dog.objects.create(name='Rufus', data={
...     'breed': 'labrador',
...     'owner': {
...         'name': 'Bob',
...         'other_pets': [{'name': 'Fishy'}],
...     },
... })
>>> Dog.objects.create(name='Meg', data={'breed': 'collie'})
```

Explain the differences between *JSON* and *JSONB*.

JSONField shares lookups with HStoreField:

- Contains (@>)
- Contained By (<@)
- Has Key (?)
- Has Any Keys (?|)
- Has Keys (?&)

## Key, Index and Path Lookups

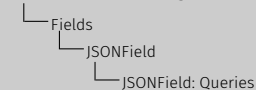
```
>>> Dog.objects.filter(data__breed='collie')
<QuerySet [<Dog: Meg>]>
```

```
>>> Dog.objects.filter(data__owner__name='Bob')
<QuerySet [<QuerySet <Dog: Rufus>]>
```

```
>>> Dog.objects.filter(data__owner__other_pets__0__name='Fishy')
<QuerySet [<Dog: Rufus>]>
```

2017-05-05

Flex schemas with PostgreSQL and Django



JSONField shares lookups with HStoreField:

- Contains (@>)
- Contained By (<@)
- Has Key (?)
- Has Any Keys (?|)
- Has Keys (?&)

### Key, Index and Path Lookups

```
>>> Dog.objects.filter(data__breed='collie')
<QuerySet [<Dog: Meg>]>

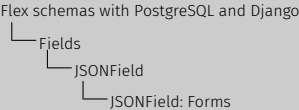
>>> Dog.objects.filter(data__owner__name='Bob')
<QuerySet [<QuerySet <Dog: Rufus>]>

>>> Dog.objects.filter(data__owner__other_pets__0__name='Fishy')
<QuerySet [<Dog: Rufus>]>
```

JSONField

A field which accepts JSON encoded data for a JSONField. It is represented by an HTML `<textarea>`.

2017-05-05



**JSONField**  
A field which accepts JSON encoded data for a JSONField. It is represented by an HTML `<textarea>`.

## Indexes

---

2017-05-05

Flex schemas with PostgreSQL and Django

└─ Indexes

Indexes

---



Hash indexes can only handle simple equality comparisons.

## Operations

= Equal to

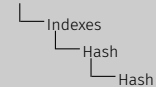
## Creation

```
from django.db import migrations

class Migration(migrations.Migration):
    operations = [
        migrations.RunSQL('CREATE INDEX IF NOT EXISTS {index_name} ON '
                           '{table_name} USING HASH ("{field_name}")'),
    ]
```

Flex schemas with PostgreSQL and Django

2017-05-05



Hash indexes can only handle simple equality comparisons.

## Operations

= Equal to

## Creation

```
from django.db import migrations

class Migration(migrations.Migration):
    operations = [
        migrations.RunSQL('CREATE INDEX IF NOT EXISTS {index_name} ON '
                           '{table_name} USING HASH ("{field_name}")'),
    ]
```

# Binary Tree (B-Tree)

B-trees can handle equality and range queries on data that can be sorted into some ordering.

B-tree indexes can also be used to retrieve data in sorted order.

## Operations

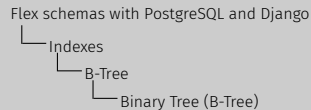
- < Less than
- <= Less than or equal to
- = Equal to
- >= Greater than or equal to
- > Greater than

## Creation

```
from django.db import models

class Foo(models.Model):
    bar = models.IntegerField(db_index=True)
```

2017-05-05



## Binary Tree (B-Tree)

B-trees can handle equality and range queries on data that can be sorted into some ordering.

B-tree indexes can also be used to retrieve data in sorted order.

### Operations

- < Less than
- <= Less than or equal to
- = Equal to
- >= Greater than or equal to
- > Greater than

### Creation

```
from django.db import models

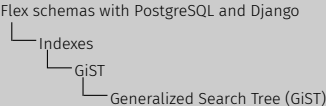
class Foo(models.Model):
    bar = models.IntegerField(db_index=True)
```

GiST indexes are not a single kind of index, it is a balanced, tree-structured access method, that acts as a base template in which to implement arbitrary indexing schemes.

## Operations

- |   |  |
|---|--|
| <b>&lt;&lt;</b> Is strictly left of?              | <b> &amp;&gt;</b> Does not extend below? |
| <b>&amp;&lt;</b> Does not extend to the right of? | <b> &gt;&gt;</b> Is strictly above?      |
| <b>&amp;&gt;</b> Does not extend to the left of?  | <b>@&gt;</b> Contains?                   |
| <b>&gt;&gt;</b> Is strictly right of?             | <b>&lt;@</b> Contained in or on?         |
| <b>&lt;&lt; </b> Is strictly below?               | <b>~=</b> Same as?                       |
| <b>&amp;&lt; </b> Does not extend above?          | <b>&amp;&amp;</b> Overlaps?              |

2017-05-05



GiST indexes are not a single kind of index, it is a balanced, tree-structured access method, that acts as a base template in which to implement arbitrary indexing schemes.

### Operations

- |   |  |
|---|--|
| <b>&lt;&lt;</b> Is strictly left of?              | <b> &amp;&gt;</b> Does not extend below? |
| <b>&amp;&lt;</b> Does not extend to the right of? | <b> &gt;&gt;</b> Is strictly above?      |
| <b>&amp;&gt;</b> Does not extend to the left of?  | <b>@&gt;</b> Contains?                   |
| <b>&gt;&gt;</b> Is strictly right of?             | <b>&lt;@</b> Contained in or on?         |
| <b>&lt;&lt; </b> Is strictly below?               | <b>~=</b> Same as?                       |
| <b>&amp;&lt; </b> Does not extend above?          | <b>&amp;&amp;</b> Overlaps?              |

# Generalized Search Tree (GiST)

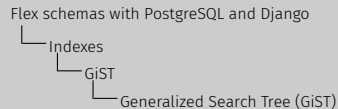
GiST indexes are not a single kind of index, it is a balanced, tree-structured access method, that acts as a base template in which to implement arbitrary indexing schemes.

## Creation

```
from django.db import migrations

class Migration(migrations.Migration):
    operations = [
        migrations.RunSQL('CREATE INDEX IF NOT EXISTS {index_name} ON '
                           '{table_name} USING GIST ("{field_name}")'),
    ]
```

2017-05-05



## Generalized Search Tree (GiST)

GiST indexes are not a single kind of index, it is a balanced, tree-structured access method, that acts as a base template in which to implement arbitrary indexing schemes.

### Creation

```
from django.db import migrations

class Migration(migrations.Migration):
    operations = [
        migrations.RunSQL('CREATE INDEX IF NOT EXISTS {index_name} ON '
                           '{table_name} USING GIST ("{field_name}")'),
    ]
```

# Space-partitioned GiST (SP-GiST)

SP-GiST supports partitioned search trees, which facilitate development of a wide range of different non-balanced data structures.

## Operations

- << Is strictly left of?
- >> Is strictly right of?
- ~= Same as?
- <@ Contained in or on?
- <^ Is below (allows touching)?
- >^ Is above (allows touching)?

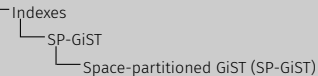
## Creation

```
from django.db import migrations

class Migration(migrations.Migration):
    operations = [
        migrations.RunSQL('CREATE INDEX IF NOT EXISTS {index_name} ON '
                           '{table_name} USING SPGIST ("{field_name}");'),
    ]
```

Flex schemas with PostgreSQL and Django

2017-05-05



## Space-partitioned GiST (SP-GiST)

SP-GiST supports partitioned search trees, which facilitate development of a wide range of different non-balanced data structures.

### Operations

- << Is strictly left of?
- >> Is strictly right of?
- ~= Same as?
- <@ Contained in or on?
- <^ Is below (allows touching)?
- >^ Is above (allows touching)?

### Creation

```
from django.db import migrations

class Migration(migrations.Migration):
    operations = [
        migrations.RunSQL('CREATE INDEX IF NOT EXISTS {index_name} ON '
                           '{table_name} USING SPGIST ("{field_name}");'),
    ]
```

# Generalized Inverted Index (GIN)

GIN is designed for handling cases where the items to be indexed are composite values, and the queries to be handled by the index need to search for element values that appear within the composite items.

## Operations

`<@` Contained in or on?

`@>` Contains?

`=` Equal to

`&&` Overlaps?

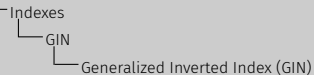
## Creation

```
from django.db import migrations
from django.contrib.postgres.indexes import GinIndex

class Migration(migrations.Migration):
    operations = [
        migrations.AddIndex("Foo", GinIndex(fields=[], name=None)),
    ]
```

Flex schemas with PostgreSQL and Django

2017-05-05



## Generalized Inverted Index (GIN)

GIN is designed for handling cases where the items to be indexed are composite values, and the queries to be handled by the index need to search for element values that appear within the composite items.

### Operations

`<@` Contained in or on?

`@>` Contains?

`=` Equal to

`&&` Overlaps?

### Creation

```
from django.db import migrations
from django.contrib.postgres.indexes import GinIndex

class Migration(migrations.Migration):
    operations = [
        migrations.AddIndex("Foo", GinIndex(fields=[], name=None)),
    ]
```

# Block Range Indexes (BRIN)

BRIN is designed for handling very large tables in which certain columns have some natural correlation with their physical location within the table.

## Operations

- < Less than
- <= Less than or equal to
- = Equal to
- >= Greater than or equal to
- > Greater than

## Creation

```
from django.db import migrations
from django.contrib.postgres.indexes import GinIndex

class Migration(migrations.Migration):
    operations = [
        migrations.AddIndex("Foo", BrinIndex(fields=[], name=None, pages_per_range=None)),
    ]
```

Flex schemas with PostgreSQL and Django

2017-05-05



## Block Range Indexes (BRIN)

BRIN is designed for handling very large tables in which certain columns have some natural correlation with their physical location within the table.

### Operations

- < Less than
- <= Less than or equal to
- = Equal to
- >= Greater than or equal to
- > Greater than

### Creation

```
from django.db import migrations
from django.contrib.postgres.indexes import GinIndex

class Migration(migrations.Migration):
    operations = [
        migrations.AddIndex("Foo", BrinIndex(fields=[], name=None, pages_per_range=None)),
    ]
```

## Conclusion

---

2017-05-05

Flex schemas with PostgreSQL and Django

└ Conclusion

Conclusion

---



# PostgreSQL as Relational and NoSQL?



Flex schemas with PostgreSQL and Django

Conclusion

PostgreSQL as Relational and NoSQL

2017-05-05

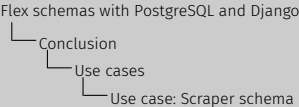
PostgreSQL as Relational and NoSQL?



Problem

Wants to scrape some websites crawling over search results and store all these results. These information must be accesible querying by website, date, and fuzzy matching over item name.

2017-05-05



Problem

Wants to scrape some websites crawling over search results and store all these results. These information must be accesible querying by website, date, and fuzzy matching over item name.

## Use case: Scraper schema

### Problem

Wants to scrape some websites crawling over search results and store all these results. These information must be accesible querying by website, date, and fuzzy matching over item name.

### Schema

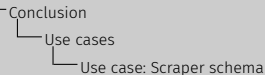
```
from django.contrib.postgres.fields import JSONField
from django.db import models
```

```
class Item(models.Model):
    website = models.URLField()
    name = models.CharField(max_length=200)
    timestamp = models.DateTimeField(auto_now_add=True)
    data = JSONField()
```

```
>>> Item.objects.annotate(
...     similarity=TrigramSimilarity('name', 'Similar item name'),
... ).filter(
...     similarity__gt=0.5,
...     timestamp__year=2017,
...     website='http://www.example.com'
... ).order_by('-similarity')
```

Flex schemas with PostgreSQL and Django

2017-05-05



Use case: Scraper schema

### Problem

Wants to scrape some websites crawling over search results and store all these results. These information must be accesible querying by website, date, and fuzzy matching over item name.

### Schema

```
from django.contrib.postgres.fields import JSONField
from django.db import models

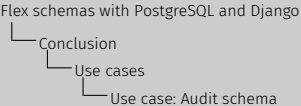
class Item(models.Model):
    website = models.URLField()
    name = models.CharField(max_length=200)
    timestamp = models.DateTimeField(auto_now_add=True)
    data = JSONField()

>>> Item.objects.annotate(
...     similarity=TrigramSimilarity('name', 'Similar item name'),
... ).filter(
...     similarity__gt=0.5,
...     timestamp__year=2017,
...     website='http://www.example.com'
... ).order_by('-similarity')
```

Problem

Integrate an audit or logging system into an application. This system generates an entry for each request-response done by users and associates it to they. Entries must contains request and response data and metadata.

2017-05-05



**Problem**  
Integrate an audit or logging system into an application. This system generates an entry for each request-response done by users and associates it to they. Entries must contains request and response data and metadata.

## Use case: Audit schema

### Problem

Integrate an audit or logging system into an application. This system generates an entry for each request-response done by users and associates it to them. Entries must contain request and response data and metadata.

### Schema

```
import datetime

from django.contrib.auth.models import User
from django.contrib.postgres.fields import JSONField
from django.db import models

class Entry(models.Model):
    user = models.ForeignKey(User)
    timestamp = models.DateTimeField(auto_now_add=True)
    request = JSONField()
    response = JSONField()

>>> Entry.objects.filter(user__id=1, timestamp=datetime.datetime.today())
```

Flex schemas with PostgreSQL and Django

2017-05-05

Conclusion  
Use cases  
Use case: Audit schema

Use case: Audit schema

### Problem

Integrate an audit or logging system into an application. This system generates an entry for each request-response done by users and associates it to them. Entries must contain request and response data and metadata.

### Schema

```
import datetime

from django.contrib.auth.models import User
from django.contrib.postgres.fields import JSONField
from django.db import models

class Entry(models.Model):
    user = models.ForeignKey(User)
    timestamp = models.DateTimeField(auto_now_add=True)
    request = JSONField()
    response = JSONField()

>>> Entry.objects.filter(user__id=1, timestamp=datetime.datetime.today())
```

# Use the right Tool for the Job



Flex schemas with PostgreSQL and Django

Conclusion

Use cases

2017-05-05

Use the right Tool for the Job

